Tidymodels

Lecture 23

Dr. Colin Rundel
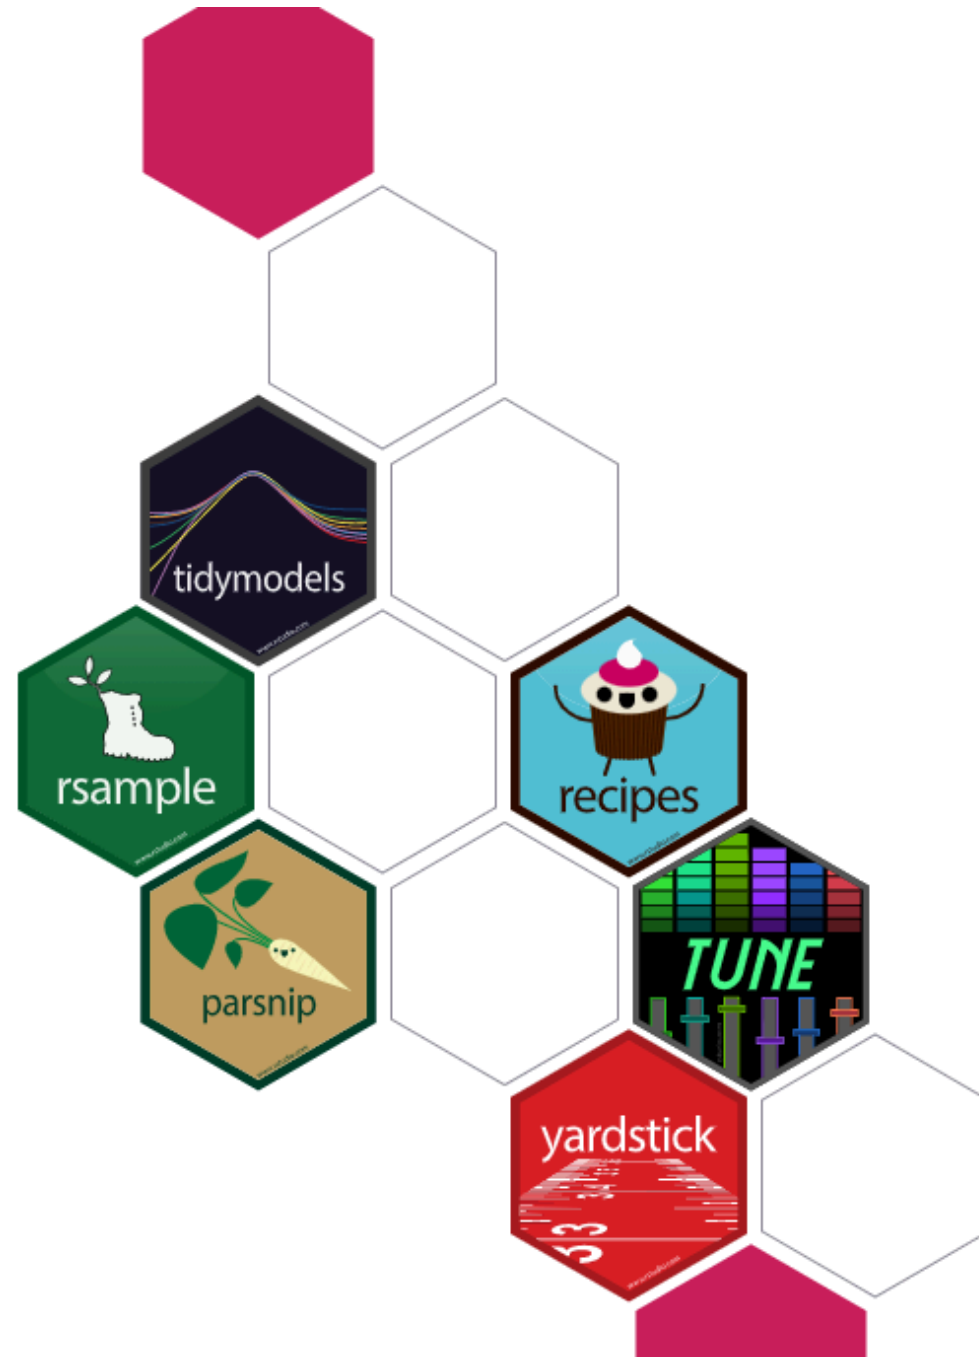
Tidymodels

```
1  library(tidymodels)
```

── Attaching packages ────────────────────────────────── tidymodels 1.3.

✔ broom	1.0.8	✔ rsample	1.3.0
✔ dials	1.4.0	✔ tune	1.3.0
✔ infer	1.0.7	✔ workflows	1.2.0
✔ modeldata	1.4.0	✔ workflowsets	1.1.0
✔ parsnip	1.3.1	✔ yardstick	1.3.2
✔ recipes	1.2.1		

── Conflicts ──────────────────────────────────── tidymodels_conflicts(

✖ scales::discard()   masks purrr::discard()
✖ dplyr::filter()   masks stats::filter()
✖ recipes::fixed()   masks stringr::fixed()
✖ dplyr::lag()   masks stats::lag()
✖ rsample::populate() masks Rcpp::populate()
✖ yardstick::spec()   masks readr::spec()
✖ recipes::step()   masks stats::step()

Book data

```r
(books = DAAG::allbacks |>
  as_tibble() |>
  select(-area) |>
  mutate(
    cover = forcats::fct_recode(
      cover,
      "hardback" = "hb",
      "paperback" = "pb"
    )
  )
)
```

```
# A tibble: 15 × 3
   volume weight cover
    <dbl>  <dbl> <fct>
 1    885    800 hardback
 2   1016    950 hardback
 3   1125   1050 hardback
 4    239    350 hardback
 5    701    750 hardback
 6    641    600 hardback
 7   1228   1075 hardback
 8    412    250 paperback
 9    953    700 paperback
10    929    650 paperback
11   1492    975 paperback
12    419    350 paperback
13   1010    950 paperback
14    595    425 paperback
15   1034    725 paperback
```

```
1  ggplot(books, aes(x=volume, y=weight, color = cover)) +
2    geom_point(size=2)
```

Building a tidymodel

```
1  linear_reg()
```

Linear Regression Model Specification (regression)

Computational engine: lm

```
1  linear_reg() |>
2    set_engine("lm")
```

Linear Regression Model Specification (regression)

Computational engine: lm

Building a tidymodel

```
1  linear_reg() |>
2    set_engine("lm") |>
3    fit(weight ~ volume * cover,
4      data = books)
```

parsnip model object

Call:
stats::lm(formula = weight ~ volume * cover, dat

Coefficients:
```
      (Intercept)                    volume
         161.58654                   0.76159
    coverpaperback   volume:coverpaperback
        -120.21407                  -0.07573
```

```
1  lm(weight ~ volume * cover, data = books)
```

Call:
lm(formula = weight ~ volume * cover, data = boo

Coefficients:
```
      (Intercept)                    volume
         161.58654                   0.76159
    coverpaperback   volume:coverpaperback
        -120.21407                  -0.07573
```

Tidy model objects

```
1  lm_tm = linear_reg() |>
2      set_engine("lm") |>
3      fit(weight ~ volume * cover,
4          data = books)
```

```
1  lm_b = lm(weight ~ volume * cover, data = bo
```

```
1  summary(lm_tm)
```

```
            Length Class      Mode
lvl          0     -none-     NULL
ordered      1     -none-     logical
spec         7     linear_reg list
fit         13     lm         list
preproc      1     -none-     list
elapsed      2     -none-     list
censor_probs 0     -none-     list
```

```
1  summary(lm_b)
```

```
Call:
lm(formula = weight ~ volume * cover, data = boo

Residuals:
   Min     1Q Median     3Q    Max
-89.67 -32.07 -21.82  17.94 215.91

Coefficients:
                         Estimate Std. Error t va
(Intercept)             161.58654   86.51918   1.
volume                    0.76159    0.09718   7.
coverpaperback         -120.21407  115.65899  -1.
volume:coverpaperback    -0.07573    0.12802  -0.
                        Pr(>|t|)
(Intercept)              0.0887 .
volume                  7.94e-06 ***
coverpaperback           0.3209
volume:coverpaperback    0.5661
---
Signif. codes:
0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
1  summary(lm_tm$fit)
```

Call:
stats::lm(formula = weight ~ volume * cover, data = data)

Residuals:
   Min   1Q Median   3Q   Max
-89.67 -32.07 -21.82 17.94 215.91

Coefficients:
                          Estimate Std. Error t value
(Intercept)   161.58654   86.51918   1.868
volume   0.76159   0.09718   7.837
coverpaperback   -120.21407 115.65899 -1.039
volume:coverpaperback   -0.07573   0.12802 -0.592
                         Pr(>|t|)
(Intercept)   0.0887 .
volume   7.94e-06 ***
coverpaperback   0.3209
volume:coverpaperback   0.5661

Tidy coefficients

```
1  broom::tidy(lm_tm)
```

```
# A tibble: 4 × 5
  term          estimate std.error statistic
  <chr>            <dbl>     <dbl>     <dbl>
1 (Intercept)   1.62e+2    86.5        1.87
2 volume        7.62e-1     0.0972     7.84
3 coverpaper…  -1.20e+2   116.        -1.04
4 volume:cov…  -7.57e-2     0.128     -0.592
```

```
1  broom::tidy(lm_b)
```

```
# A tibble: 4 × 5
  term          estimate std.error statistic
  <chr>            <dbl>     <dbl>     <dbl>
1 (Intercept)   1.62e+2    86.5        1.87
2 volume        7.62e-1     0.0972     7.84
3 coverpaper…  -1.20e+2   116.        -1.04
4 volume:cov…  -7.57e-2     0.128     -0.592
```

Tidy statistics

```
1 broom::glance(lm_tm)
```

```
# A tibble: 1 × 12
  r.squared adj.r.squared sigma statistic  p.value    df logLik   AIC   BIC deviance
      <dbl>         <dbl> <dbl>     <dbl>    <dbl> <dbl>  <dbl> <dbl> <dbl>    <dbl>
1     0.930         0.911  80.4      48.5  1.24e-6     3  -84.8  180.  183.   71118.
# ℹ 2 more variables: df.residual <int>, nobs <int>
```

```
1 broom::glance(lm_b)
```

```
# A tibble: 1 × 12
  r.squared adj.r.squared sigma statistic  p.value    df logLik   AIC   BIC deviance
      <dbl>         <dbl> <dbl>     <dbl>    <dbl> <dbl>  <dbl> <dbl> <dbl>    <dbl>
1     0.930         0.911  80.4      48.5  1.24e-6     3  -84.8  180.  183.   71118.
# ℹ 2 more variables: df.residual <int>, nobs <int>
```

Tidy prediction

```
1  broom::augment(lm_tm, new_data = books)
```

```
# A tibble: 15 × 5
   .pred .resid volume weight cover
   <dbl>  <dbl>  <dbl>  <dbl> <fct>
 1  836. -35.6     885    800 hardback
 2  935.  14.6    1016    950 hardback
 3 1018.  31.6    1125   1050 hardback
 4  344.   6.39    239    350 hardback
 5  695.  54.5     701    750 hardback
 6  650. -49.8     641    600 hardback
 7 1097. -21.8    1228   1075 hardback
 8  324. -73.9     412    250 paperback
 9  695.   5.00    953    700 paperback
10  679. -28.5     929    650 paperback
11 1065. -89.7    1492    975 paperback
12  329.  21.3     419    350 paperback
13  734. 216.     1010    950 paperback
14  449. -24.5     595    425 paperback
15  751. -25.6    1034    725 paperback
```

Putting it together

```r
lm_tm |>
  augment(
    new_data = tidyr::expand_grid(
      volume = seq(0, 1500, by=5),
      cover = c("hardback", "paperback") |> as.factor()
    )
  ) |>
  rename(weight = .pred) |>
  ggplot(aes(x = volume, y = weight, color = cover, group = cover)) +
    geom_line() +
    geom_point(data = books)
```

parsnip

www.rstudio.com

Why do we care?

```
1  show_engines("linear_reg")
```

```
# A tibble: 8 × 2
  engine    mode
  <chr>     <chr>
1 lm        regression
2 glm       regression
3 glmnet    regression
4 stan      regression
5 spark     regression
6 keras     regression
7 brulee    regression
8 quantreg  quantile regression
```

```
1  (bayes_tm = linear_reg() |>
2    set_engine(
3      "stan",
4      prior_intercept = rstanarm::student_t(d
5      prior = rstanarm::student_t(df = 1),
6      seed = 1234
7    )
8  )
```

```
Linear Regression Model Specification (regressio

Engine-Specific Arguments:
  prior_intercept = rstanarm::student_t(df = 1)
  prior = rstanarm::student_t(df = 1)
  seed = 1234

Computational engine: stan
```

Fitting with `rstanarm`

```
1  (bayes_tm = bayes_tm |>
2    fit(weight ~ volume * cover, data = books)
3  )
```

Warning: There were 36 divergent transitions after warmup. See
https://mc-stan.org/misc/warnings.html#divergent-transitions-after-warmup
to find out why this is a problem and how to eliminate them.

Warning: Examine the pairs() plot to diagnose sampling problems

parsnip model object

stan_glm
 family: gaussian [identity]
 formula: weight ~ volume * cover
 observations: 15
 predictors: 4

                       Median MAD_SD
(Intercept)   94.9   60.4
volume   0.8   0.1
coverpaperback   -0.4   3.6
volume:coverpaperback   -0.2   0.0

Auxiliary parameter(s):
      Median MAD_SD
sigma 84.9   17.6

What was actually run?

```r
1  linear_reg() |>
2    set_engine(
3      "stan",
4      prior_intercept = rstanarm::student_t(df = 1),
5      prior = rstanarm::student_t(df = 1),
6      seed = 1234
7    ) |>
8    translate()
```

```
Linear Regression Model Specification (regression)

Engine-Specific Arguments:
  prior_intercept = rstanarm::student_t(df = 1)
  prior = rstanarm::student_t(df = 1)
  seed = 1234

Computational engine: stan

Model fit template:
rstanarm::stan_glm(formula = missing_arg(), data = missing_arg(),
    weights = missing_arg(), prior_intercept = rstanarm::student_t(df = 1),
    prior = rstanarm::student_t(df = 1), seed = 1234, family = stats::gaussian,
    refresh = 0)
```

Back to broom

```
1  broom::tidy(bayes_tm)
```

Error in `generics::tidy()`:
! `x` seems to be outputted from the rstanarm package.
ℹ Tidiers for mixed model output now live in broom.mixed.

```
1  broom.mixed::tidy(bayes_tm)
```

```
# A tibble: 4 × 3
  term                  estimate std.error
  <chr>                    <dbl>    <dbl>
1 (Intercept)              94.9     60.4
2 volume                    0.834    0.0734
3 coverpaperback           -0.354    3.58
4 volume:coverpaperback    -0.199    0.0498
```

```
1  broom.mixed::glance(bayes_tm)
```

```
# A tibble: 1 × 4
  algorithm   pss  nobs sigma
  <chr>     <dbl> <int> <dbl>
1 sampling   4000    15  84.9
```

Augment

```
1  augment(bayes_tm, new_data=books)
```

```
# A tibble: 15 × 5
    .pred   .resid volume weight cover
    <dbl>    <dbl>  <dbl>  <dbl> <fct>
 1  831.    -30.7     885    800 hardback
 2  940.     10.3    1016    950 hardback
 3 1030.     19.7    1125   1050 hardback
 4  293.     56.5     239    350 hardback
 5  678.     72.3     701    750 hardback
 6  628.    -27.8     641    600 hardback
 7 1116.    -41.0    1228   1075 hardback
 8  354.   -104.      412    250 paperback
 9  696.      4.18    953    700 paperback
10  681.    -30.6     929    650 paperback
11 1037.    -61.8    1492    975 paperback
12  358.     -8.00    419    350 paperback
13  732.    218.     1010    950 paperback
14  469.    -44.3     595    425 paperback
15  747.    -22.1    1034    725 paperback
```

Predictions

```
1  bayes_tm |>
2    augment(
3      new_data = tidyr::expand_grid(
4        volume = seq(0, 1500, by=5),
5        cover = c("hardback", "paperback") |> as.factor()
6      )
7    ) |>
8    rename(weight = .pred) |>
9    ggplot(aes(x = volume, y = weight, color = cover, group = cover)) +
10     geom_line() +
11     geom_point(data = books)
```

Performance

```
1  lm_tm |>
2    augment(new_data = books) |>
3    yardstick::rmse(weight, .pred)
```

```
# A tibble: 1 × 3
  .metric .estimator .estimate
  <chr>   <chr>          <dbl>
1 rmse    standard        68.9
```

```
1  bayes_tm |>
2    augment(new_data = books) |>
3    yardstick::rmse(weight, .pred)
```

```
# A tibble: 1 × 3
  .metric .estimator .estimate
  <chr>   <chr>          <dbl>
1 rmse    standard        72.1
```

Cross validation and Feature engineering

The Office & IMDB

```
1 (office_ratings = read_csv("data/office_ratings.csv"))
```

```
# A tibble: 188 × 6
   season episode title              imdb_rating total_votes air_date
    <dbl>   <dbl> <chr>                    <dbl>       <dbl> <date>
 1      1       1 Pilot                      7.6        3706 2005-03-24
 2      1       2 Diversity Day              8.3        3566 2005-03-29
 3      1       3 Health Care                7.9        2983 2005-04-05
 4      1       4 The Alliance               8.1        2886 2005-04-12
 5      1       5 Basketball                 8.4        3179 2005-04-19
 6      1       6 Hot Girl                   7.8        2852 2005-04-26
 7      2       1 The Dundies                8.7        3213 2005-09-20
 8      2       2 Sexual Harassment          8.2        2736 2005-09-27
 9      2       3 Office Olympics            8.4        2742 2005-10-04
10      2       4 The Fire                   8.4        2713 2005-10-11
# ℹ 178 more rows
```

Rating vs Air Date

Test-train split

```
1  set.seed(123)
2  (office_split = initial_split(office_ratings, prop = 0.8))
```

```
<Training/Testing/Total>
<150/38/188>
```

```
1  (office_train = training(office_split))
```

```
# A tibble: 150 × 6
   season episode title    imdb_rating total_vot
    <dbl>   <dbl> <chr>           <dbl>      <db
 1      8      18 Last Da…          7.8         14
 2      9      14 Vandali…          7.6         14
 3      2       8 Perform…          8.2         24
 4      9       5 Here Co…          7.1         15
 5      3      22 Beach G…          9.1         27
 6      7       1 Nepotism          8.4         18
 7      3      15 Phyllis…          8.3         22
 8      9      21 Livin' …          8.9         20
 9      9      18 Promos            8           14
10      8      12 Pool Pa…          8           16
# ℹ 140 more rows
# ℹ 1 more variable: air_date <date>
```

```
1  (office_test = testing(office_split))
```

```
# A tibble: 38 × 6
   season episode title    imdb_rating total_vot
    <dbl>   <dbl> <chr>           <dbl>      <db
 1      1       2 Diversi…          8.3         35
 2      2       4 The Fire          8.4         27
 3      2       9 E-Mail …          8.4         25
 4      2      12 The Inj…          9           32
 5      2      22 Casino …          9.3         36
 6      3       5 Initiat…          8.2         22
 7      3      16 Busines…          8.8         26
 8      3      17 Cocktai…          8.5         22
 9      4       6 Branch …          8.5         21
10      4       7 Survivo…          8.3         21
# ℹ 28 more rows
# ℹ 1 more variable: air_date <date>
```

Feature engineering with dplyr

```r
1  office_train |>
2    mutate(
3      season = as_factor(season),
4      month = lubridate::month(air_date),
5      wday = lubridate::wday(air_date),
6      top10_votes = as.integer(total_votes > quantile(total_votes, 0.9))
7    )
```

```
# A tibble: 150 × 9
   season episode title                imdb_rating total_votes air_date    month  wday top10_votes
   <fct>    <dbl> <chr>                        <dbl>       <dbl> <date>      <dbl> <dbl>       <int>
 1 8           18 Last Day in Florida           7.8        1429 2012-03-08      3     5           0
 2 9           14 Vandalism                     7.6        1402 2013-01-31      1     5           0
 3 2            8 Performance Review            8.2        2416 2005-11-15     11     3           0
 4 9            5 Here Comes Treble             7.1        1515 2012-10-25     10     5           0
 5 3           22 Beach Games                   9.1        2783 2007-05-10      5     5           0
 6 7            1 Nepotism                      8.4        1897 2010-09-23      9     5           0
 7 3           15 Phyllis' Wedding              8.3        2283 2007-02-08      2     5           0
 8 9           21 Livin' the Dream              8.9        2041 2013-05-02      5     5           0
 9 9           18 Promos                        8          1445 2013-04-04      4     5           0
10 8           12 Pool Party                    8          1612 2012-01-19      1     5           0
# ℹ 140 more rows
```

Anyone see a potential problem with the code above?

Better living through recipes

```r
1  r = recipe(imdb_rating ~ ., data = office_train)
```

```r
1  summary(r)
```

```
# A tibble: 6 × 4
  variable     type      role       source
  <chr>        <list>    <chr>      <chr>
1 season       <chr [2]> predictor original
2 episode      <chr [2]> predictor original
3 title        <chr [3]> predictor original
4 total_votes  <chr [2]> predictor original
5 air_date     <chr [1]> predictor original
6 imdb_rating  <chr [2]> outcome   original
```

Recipe roles

```r
1  r = recipe(
2    imdb_rating ~ ., data = office_train
3  ) |>
4    update_role(title, new_role = "ID")
```

```r
1  summary(r)
```

```
# A tibble: 6 × 4
  variable    type       role      source
  <chr>       <list>     <chr>     <chr>
1 season      <chr [2]>  predictor original
2 episode     <chr [2]>  predictor original
3 title       <chr [3]>  ID        original
4 total_votes <chr [2]>  predictor original
5 air_date    <chr [1]>  predictor original
6 imdb_rating <chr [2]>  outcome   original
```

Adding features (month & day of week)

```r
1  r = recipe(
2    imdb_rating ~ ., data = office_train
3  ) |>
4    update_role(title, new_role = "ID") |>
5    step_date(air_date, features = c("dow", "month"))
```

Adding Holidays

```
 1  r = recipe(
 2    imdb_rating ~ ., data = office_train
 3  ) |>
 4    update_role(title, new_role = "ID") |>
 5    step_date(air_date, features = c("dow", "month")) |>
 6    step_holiday(
 7      air_date,
 8      holidays = c("USThanksgivingDay", "USChristmasDay", "USNewYearsDay", "USIndepe
 9      keep_original_cols = FALSE
10    )
```

Seasons as factors

```r
 1  r = recipe(
 2    imdb_rating ~ ., data = office_train
 3  ) |>
 4    update_role(title, new_role = "ID") |>
 5    step_date(air_date, features = c("dow", "month")) |>
 6    step_holiday(
 7      air_date,
 8      holidays = c("USThanksgivingDay", "USChristmasDay", "USNewYearsDay", "USIndepe
 9      keep_original_cols = FALSE
10    ) |>
11    step_num2factor(season, levels = as.character(1:9))
```

Dummy coding

```r
 1  r = recipe(
 2    imdb_rating ~ ., data = office_train
 3  ) |>
 4    update_role(title, new_role = "ID") |>
 5    step_date(air_date, features = c("dow", "month")) |>
 6    step_holiday(
 7      air_date,
 8      holidays = c("USThanksgivingDay", "USChristmasDay", "USNewYearsDay", "USIndepe
 9      keep_original_cols = FALSE
10    ) |>
11    step_num2factor(season, levels = as.character(1:9)) |>
12    step_dummy(all_nominal_predictors())
```

top10_votes

```r
1  r = recipe(
2    imdb_rating ~ ., data = office_train
3  ) |>
4    update_role(title, new_role = "ID") |>
5    step_date(air_date, features = c("dow", "month")) |>
6    step_holiday(
7      air_date,
8      holidays = c("USThanksgivingDay", "USChristmasDay", "USNewYearsDay", "USIndepe
9      keep_original_cols = FALSE
10   ) |>
11   step_num2factor(season, levels = as.character(1:9)) |>
12   step_dummy(all_nominal_predictors()) |>
13   step_percentile(total_votes) |>
14   step_mutate(top10 = as.integer(total_votes >= 0.9)) |>
15   step_rm(total_votes)
```

Preparing a recipe

```
1  prep(r)
```

── Recipe ──────────────────────────────

── Inputs

Number of variables by role

outcome: 1
predictor: 4
ID: 1

── Training information

Training data contained 150 data points and no incomplete rows.

── Operations

• Date features from: air_date | Trained

• Holiday features from: air_date | Trained

• Factor variables from: season | Trained

• Dummy variables from: season, air_date_dow, air_date_month | Trained

• Percentile transformation on: total_votes | Trained

• Variable mutation for: ~as.integer(total_votes >= 0.9) | Trained

• Variables removed: total_votes | Trained

Baking a recipe

```
1  prep(r) |>
2    bake(new_data = office_train)
```

```
# A tibble: 150 × 33
   episode title     imdb_rating air_date_USThanksgiv…¹ air_date_USChristmas…² air_date_USNewYears[
     <dbl> <chr>           <dbl>                  <int>                  <int>                   <in
 1      18 Last Da…         7.8                       0                      0                      0
 2      14 Vandali…         7.6                       0                      0                      0
 3       8 Perform…         8.2                       0                      0                      0
 4       5 Here Co…         7.1                       0                      0                      0
 5      22 Beach G…         9.1                       0                      0                      0
 6       1 Nepotism         8.4                       0                      0                      0
 7      15 Phyllis…         8.3                       0                      0                      0
 8      21 Livin' …         8.9                       0                      0                      0
 9      18 Promos          8                         0                      0                      0
10      12 Pool Pa…         8                         0                      0                      0
# ℹ 140 more rows
# ℹ abbreviated names: ¹air_date_USThanksgivingDay, ²air_date_USChristmasDay
# ℹ 27 more variables: air_date_USIndependenceDay <int>, season_X2 <dbl>, season_X3 <dbl>,
#   season_X4 <dbl>, season_X5 <dbl>, season_X6 <dbl>, season_X7 <dbl>, season_X8 <dbl>,
#   season_X9 <dbl>, air_date_dow_Mon <dbl>, air_date_dow_Tue <dbl>, air_date_dow_Wed <dbl>,
#   air_date_dow_Thu <dbl>, air_date_dow_Fri <dbl>, air_date_dow_Sat <dbl>,
#   air_date_month_Feb <dbl>, air_date_month_Mar <dbl>, air_date_month_Apr <dbl>, …
```

Informative features?

```
1  prep(r) |>
2    bake(new_data = office_train) |>
3    map_int(~ length(unique(.x)))
```

episode	title	imdb_rating
26	150	26
air_date_USThanksgivingDay	air_date_USChristmasDay	air_date_USNewYearsDay
1	1	1
air_date_USIndependenceDay	season_X2	season_X3
1	2	2
season_X4	season_X5	season_X6
2	2	2
season_X7	season_X8	season_X9
2	2	2
air_date_dow_Mon	air_date_dow_Tue	air_date_dow_Wed
1	2	1
air_date_dow_Thu	air_date_dow_Fri	air_date_dow_Sat
2	1	1
air_date_month_Feb	air_date_month_Mar	air_date_month_Apr
2	2	2
air_date_month_May	air_date_month_Jun	air_date_month_Jul
2	1	1
air_date_month_Aug	air_date_month_Sep	air_date_month_Oct
1	2	2
air_date_month_Nov	air_date_month_Dec	top10
2	2	2

Removing zero variance predictors

```r
1  r = recipe(
2      imdb_rating ~ ., data = office_train
3    ) |>
4    update_role(title, new_role = "ID") |>
5    step_date(air_date, features = c("dow", "month")) |>
6    step_holiday(
7      air_date,
8      holidays = c("USThanksgivingDay", "USChristmasDay", "USNewYearsDay", "USIndepe
9      keep_original_cols = FALSE
10   ) |>
11   step_num2factor(season, levels = as.character(1:9)) |>
12   step_dummy(all_nominal_predictors()) |>
13   step_percentile(total_votes) |>
14   step_mutate(top10 = as.integer(total_votes >= 0.9)) |>
15   step_rm(total_votes) |>
16   step_zv(all_predictors())
```

```
1  prep(r) |>
2    bake(new_data = office_train)
```

```
# A tibble: 150 × 22
   episode title   imdb_rating season_X2 season_X3 season_X4 season_X5 season_X6 season_X7 season_
     <dbl> <chr>         <dbl>     <dbl>     <dbl>     <dbl>     <dbl>     <dbl>     <dbl>      <dk
 1      18 Last D…         7.8         0         0         0         0         0         0
 2      14 Vandal…         7.6         0         0         0         0         0         0
 3       8 Perfor…         8.2         1         0         0         0         0         0
 4       5 Here C…         7.1         0         0         0         0         0         0
 5      22 Beach …         9.1         0         1         0         0         0         0
 6       1 Nepoti…         8.4         0         0         0         0         0         1
 7      15 Phylli…         8.3         0         1         0         0         0         0
 8      21 Livin'…         8.9         0         0         0         0         0         0
 9      18 Promos          8           0         0         0         0         0         0
10      12 Pool P…         8           0         0         0         0         0         0
# ℹ 140 more rows
# ℹ 12 more variables: season_X9 <dbl>, air_date_dow_Tue <dbl>, air_date_dow_Thu <dbl>,
#   air_date_month_Feb <dbl>, air_date_month_Mar <dbl>, air_date_month_Apr <dbl>,
#   air_date_month_May <dbl>, air_date_month_Sep <dbl>, air_date_month_Oct <dbl>,
#   air_date_month_Nov <dbl>, air_date_month_Dec <dbl>, top10 <int>
```

Sta 323 - Spring 2025

43

Really putting it all together

```r
(office_work = workflow() |>
  add_recipe(r) |>
  add_model(
    linear_reg() |>
    set_engine("lm")
  )
)
```

```
══ Workflow ════════════════════════════════════════════
Preprocessor: Recipe
Model: linear_reg()

── Preprocessor ────────────────────────────────────────
8 Recipe Steps

• step_date()
• step_holiday()
• step_num2factor()
• step_dummy()
• step_percentile()
• step_mutate()
• step_rm()
• step_zv()

── Model ───────────────────────────────────────────────
Linear Regression Model Specification (regression)

Computational engine: lm
```

Workflow fit

```
1  (office_fit = office_work |>
2    fit(data = office_train))
```

```
══ Workflow [trained] ══════════════════════════════════
Preprocessor: Recipe
Model: linear_reg()

── Preprocessor ─────────────────────────────────────────
8 Recipe Steps

• step_date()
• step_holiday()
• step_num2factor()
• step_dummy()
• step_percentile()
• step_mutate()
• step_rm()
• step_zv()

── Model ────────────────────────────────────────────────

Call:
stats::lm(formula = ..y ~ ., data = data)

Coefficients:
    (Intercept)          episode        season_X2          season_X3           season_X4
```

Performance

```
1  office_fit |>
2    augment(office_train) |>
3    rmse(imdb_rating, .pred)
```

```
# A tibble: 1 × 3
  .metric .estimator .estimate
  <chr>   <chr>          <dbl>
1 rmse    standard       0.342
```

```
1  office_fit |>
2    augment(office_test) |>
3    rmse(imdb_rating, .pred)
```

```
# A tibble: 1 × 3
  .metric .estimator .estimate
  <chr>   <chr>          <dbl>
1 rmse    standard       0.399
```

k-fold cross validation

Creating folds

```
set.seed(123)
(folds = vfold_cv(office_train, v=5))
```

```
#  5-fold cross-validation
# A tibble: 5 × 2
  splits            id
  <list>            <chr>
1 <split [120/30]> Fold1
2 <split [120/30]> Fold2
3 <split [120/30]> Fold3
4 <split [120/30]> Fold4
5 <split [120/30]> Fold5
```

```
(office_fit_folds = office_work |>
   fit_resamples(folds)
)
```

```
→ A | warning: prediction from rank-defici

There were issues with some computations
There were issues with some computations

# Resampling results
# 5-fold cross-validation
# A tibble: 5 × 4
  splits            id    .metrics
  <list>            <chr> <list>
1 <split [120/30]> Fold1 <tibble [2 × 4]>
2 <split [120/30]> Fold2 <tibble [2 × 4]>
3 <split [120/30]> Fold3 <tibble [2 × 4]>
4 <split [120/30]> Fold4 <tibble [2 × 4]>
5 <split [120/30]> Fold5 <tibble [2 × 4]>
  .notes
  <list>
1 <tibble [0 × 3]>
2 <tibble [1 × 3]>
3 <tibble [0 × 3]>
4 <tibble [0 × 3]>
5 <tibble [0 × 3]>

There were issues with some computations:
```

Fold performance

```
1 tune::collect_metrics(office_fit_folds)
```

```
# A tibble: 2 × 6
  .metric .estimator  mean     n std_err .config
  <chr>   <chr>      <dbl> <int>   <dbl> <chr>
1 rmse    standard   0.420     5  0.0182 Preprocessor1_Model1
2 rsq     standard   0.429     5  0.0597 Preprocessor1_Model1
```

```
1 tune::collect_metrics(office_fit_folds, summarize = FALSE) |>
2   filter(.metric == "rmse")
```

```
# A tibble: 5 × 5
  id    .metric .estimator .estimate .config
  <chr> <chr>   <chr>          <dbl> <chr>
1 Fold1 rmse    standard       0.467 Preprocessor1_Model1
2 Fold2 rmse    standard       0.403 Preprocessor1_Model1
3 Fold3 rmse    standard       0.405 Preprocessor1_Model1
4 Fold4 rmse    standard       0.454 Preprocessor1_Model1
5 Fold5 rmse    standard       0.368 Preprocessor1_Model1
```

More on the tune package next time