

Profiling & Parallelization

Lecture 21

Dr. Colin Rundel

Profiling & Benchmarking

profvis demo

```
1 n = 1e6
2 d = tibble(
3   x1 = rt(n, df = 3),
4   x2 = rt(n, df = 3),
5   x3 = rt(n, df = 3),
6   x4 = rt(n, df = 3),
7   x5 = rt(n, df = 3),
8 ) |>
9   mutate(y = -2*x1 - 1*x2 + 0*x3 + 1*x4 + 2*x5 + rnorm(n))
```

```
1 profvis::profvis({
2   lm(y~., data=d)
3 })
```

profvis demo 2

```
1 profvis::profvis({
2   data = data.frame(value = runif(5e4))
3
4   data$sum[1] = data$value[1]
5   for (i in seq(2, nrow(data))) {
6     data$sum[i] = data$sum[i-1] + data$value[i]
7   }
8 })
```

```
1 profvis::profvis({
2   x = runif(5e4)
3   sum = x[1]
4   for (i in seq(2, length(x))) {
5     sum[i] = sum[i-1] + x[i]
6   }
7 })
```

Benchmarking - bench

```
1 d = tibble(  
2   x = runif(10000),  
3   y = runif(10000)  
4 )  
5  
6 (b = bench::mark(  
7   d[d$x > 0.5, ],  
8   d[which(d$x > 0.5), ],  
9   subset(d, x > 0.5),  
10  filter(d, x > 0.5)  
11 ))
```

A tibble: 4 × 6

expression	min	median	`itr/sec`	mem_alloc	`gc/sec`
<bch:expr>	<bch:tm>	<bch:tm>	<dbl>	<bch:byt>	<dbl>
1 d[d\$x > 0.5,]	39.3µs	43.9µs	21703.	235.44KB	60.4
2 d[which(d\$x > 0.5),]	55.1µs	58.9µs	16690.	269.94KB	86.5
3 subset(d, x > 0.5)	63.5µs	67.9µs	14549.	288.01KB	73.8
4 filter(d, x > 0.5)	303.9µs	327.3µs	2960.	1.47MB	16.7

Larger n

```
1 d = tibble(  
2   x = runif(1e6),  
3   y = runif(1e6)  
4 )  
5  
6 (b = bench::mark(  
7   d[d$x > 0.5, ],  
8   d[which(d$x > 0.5), ],  
9   subset(d, x > 0.5),  
10  filter(d, x > 0.5)  
11 ))
```

A tibble: 4 × 6

expression	min	median	`itr/sec`	mem_alloc	`gc/sec`
<bch:expr>	<bch:tm>	<bch:tm>	<dbl>	<bch:byt>	<dbl>
1 d[d\$x > 0.5,]	2.81ms	2.9ms	315.	13.4MB	242.
2 d[which(d\$x > 0.5),]	6.03ms	6.18ms	147.	24.8MB	264.
3 subset(d, x > 0.5)	6.75ms	6.91ms	129.	24.8MB	249.
4 filter(d, x > 0.5)	4.13ms	4.38ms	186.	24.8MB	299.

bench - relative results

```
1 summary(b, relative=TRUE)
```

```
# A tibble: 4 × 6
```

	expression	min	median	`itr/sec`	mem_alloc	`gc/sec`
	<bch:expr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	d[d\$x > 0.5,]	1	1	2.45	1	1
2	d[which(d\$x > 0.5),]	2.15	2.13	1.14	1.86	1.09
3	subset(d, x > 0.5)	2.41	2.38	1	1.86	1.03
4	filter(d, x > 0.5)	1.47	1.51	1.45	1.86	1.24

t.test

Imagine we have run 1000 experiments (rows), each of which collects data on 50 individuals (columns).
The first 25 individuals in each experiment are assigned to group 1 and the rest to group 2.

The goal is to calculate the t-statistic for each experiment comparing group 1 to group 2.

```
1 m = 1000
2 n = 50
3 X = matrix(
4   rnorm(m * n, mean = 10, sd = 3),
5   ncol = m
6 ) |>
7 as.data.frame() |>
8 set_names(paste0("exp", seq_len(m))) |>
9 mutate(
10   ind = seq_len(n),
11   group = rep(1:2, each = n/2)
12 ) |>
13 as_tibble() |>
14 relocate(ind, group)
```

```
1 X
```

```
# A tibble: 50 × 1,002
      ind group  exp1  exp2  exp3  exp4  exp5  exp6
  <int> <int> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1     1     1     8.66 10.5   5.83  5.35 10.9   12.8
2     2     1    17.0 13.8  14.2  5.30  7.47 10.5
3     3     1     6.54  9.09 14.3  10.7 11.4   8.26
4     4     1     7.04 11.7  14.2  5.76 10.1   1.50
5     5     1     8.41  5.29  8.96 13.1   8.49  7.66
6     6     1    11.0  6.49  6.44  9.23  9.80  9.62
7     7     1     8.15 10.2  11.6 12.0   6.85  8.03
8     8     1    11.4 12.3   9.82  8.48 11.2   8.08
9     9     1    12.8  9.49 12.6 17.3   8.99  9.11
10    10     1     9.60  6.23 11.7 11.8   8.27 11.0
# i 40 more rows
# i 994 more variables: exp7 <dbl>, exp8 <dbl>,
#   exp9 <dbl>, exp10 <dbl>, exp11 <dbl>,
#   exp12 <dbl>, exp13 <dbl>, exp14 <dbl>,
#   exp15 <dbl>, exp16 <dbl>, exp17 <dbl>,
#   exp18 <dbl>, exp19 <dbl>, exp20 <dbl>,
```


Implementations

```
1 ttest_formula = function(X, m) {  
2   for(i in 1:m) t.test(X[[2+i]] ~ X$group)$stat  
3 }  
4 system.time(ttest_formula(X,m))
```

```
user  system elapsed  
0.197   0.000   0.198
```

```
1 ttest_for = function(X, m) {  
2   for(i in 1:m) t.test(X[[2+i]][X$group == 1], X[[2+i]][X$group == 2])$stat  
3 }  
4 system.time(ttest_for(X,m))
```

```
user  system elapsed  
0.063   0.000   0.063
```

```
1 ttest_apply = function(X) {  
2   f = function(x, g) {  
3     t.test(x[g==1], x[g==2])$stat  
4   }  
5   apply(X[,-(1:2)], 2, f, X$group)  
6 }  
7 system.time(ttest_apply(X))
```

```
user  system elapsed  
0.054   0.000   0.055
```

Implementations (cont.)

```
1 ttest_hand_calc = function(X) {  
2   f = function(x, grp) {  
3     t_stat = function(x) {  
4       m = mean(x)  
5       n = length(x)  
6       var = sum((x - m) ^ 2) / (n - 1)  
7  
8       list(m = m, n = n, var = var)  
9     }  
10  
11    g1 = t_stat(x[grp == 1])  
12    g2 = t_stat(x[grp == 2])  
13  
14    se_total = sqrt(g1$var / g1$n + g2$var / g2$n)  
15    (g1$m - g2$m) / se_total  
16  }  
17  
18  apply(X[,-(1:2)], 2, f, X$group)  
19 }  
20 system.time(ttest_hand_calc(X))
```

user	system	elapsed
0.016	0.000	0.016

Comparison

```
1 bench::mark(  
2   ttest_formula(X, m),  
3   ttest_for(X, m),  
4   ttest_apply(X),  
5   ttest_hand_calc(X),  
6   check=FALSE  
7 )
```

Warning: Some expressions had a GC in every iteration; so filtering is disabled.

A tibble: 4 × 6

	expression	min	median	`itr/sec`	mem_alloc	`gc/sec`
	<bch:expr>	<bch:tm>	<bch:tm>	<dbl>	<bch:byt>	<dbl>
1	ttest_formula(X, m)	218.73ms	224.1ms	4.47	8.24MB	23.8
2	ttest_for(X, m)	69.59ms	73.2ms	13.7	1.91MB	25.4
3	ttest_apply(X)	63.28ms	68.4ms	14.7	3.48MB	25.8
4	ttest_hand_calc(X)	9.36ms	10.1ms	72.4	3.44MB	21.5

Parallelization

parallel

Part of the base packages in R

- tools for the forking of R processes (some functions do not work on Windows)
- Core functions:
 - `detectCores`
 - `pvec`
 - `mclapply`
 - `mcpipeline` & `mccollect`

detectCores

Surprisingly, detects the number of cores of the current system.

```
1 detectCores()
```

```
[1] 32
```

pvec

Parallelization of a vectorized function call

```
1 system.time(pvec(1:1e7, sqrt, mc.cores = 1))
```

user	system	elapsed
0.028	0.038	0.068

```
1 system.time(pvec(1:1e7, sqrt, mc.cores = 4))
```

user	system	elapsed
0.211	0.253	0.324

```
1 system.time(pvec(1:1e7, sqrt, mc.cores = 8))
```

user	system	elapsed
0.101	0.288	0.176

```
1 system.time(sqrt(1:1e7))
```

user	system	elapsed
0.024	0.042	0.066

pvec - bench::system_time

```
1 bench::system_time(pvec(1:1e7, sqrt, mc.cores = 1))
```

process	real
125ms	126ms

```
1 bench::system_time(pvec(1:1e7, sqrt, mc.cores = 4))
```

process	real
220ms	250ms

```
1 bench::system_time(pvec(1:1e7, sqrt, mc.cores = 8))
```

process	real
272ms	286ms


```
1 bench::system_time(Sys.sleep(.5))
```

```
process    real  
54.4µs 500.8ms
```

```
1 system.time(Sys.sleep(.5))
```

```
   user  system elapsed  
0.000   0.000   0.501
```

Cores by size

```
1 cores = c(1,4,6,8,10)
2 order = 6:8
3 f = function(x,y) {
4   system.time(
5     pvec(1:(10^y), sqrt, mc.cores = x)
6   )[3]
7 }
8
9 res = map(
10  cores,
11  function(x) {
12    map_dbl(order, f, x = x)
13  }
14 ) |>
15 do.call(rbind, args = _)
16
17 rownames(res) = paste0(cores," cores")
18 colnames(res) = paste0("10^",order)
```

```
1 res
```

		10 ⁶	10 ⁷	10 ⁸
1 cores		0.004	0.126	0.745
4 cores		0.033	0.171	2.190
6 cores		0.038	0.173	1.847
8 cores		0.045	0.183	1.827
10 cores		0.050	0.187	1.852

mclapply

implements a parallelized version of `lapply`

```
1 system.time(rnorm(1e7))
```

```
user  system elapsed
0.197   0.020   0.219
```

```
1 system.time(unlist(mclapply(1:10, function(x) rnorm(1e6), mc.cores = 2)))
```

```
user  system elapsed
0.247   0.200   0.260
```

```
1 system.time(unlist(mclapply(1:10, function(x) rnorm(1e6), mc.cores = 4)))
```

```
user  system elapsed
0.231   0.216   0.172
```

```
1 system.time(unlist(mclapply(1:10, function(x) rnorm(1e6), mc.cores = 8)))
```

```
user  system elapsed
0.238   0.271   0.149
```

```
1 system.time(unlist(mclapply(1:10, function(x) rnorm(1e6), mc.cores = 10)))
```

```
user  system elapsed
0.246   0.332   0.144
```

mcparallel

Asynchronously evaluation of an R expression in a separate process

```
1 m = mcparallel(rnorm(1e6))
2 n = mcparallel(rbeta(1e6,1,1))
3 o = mcparallel(rgamma(1e6,1,1))
```

```
1 str(m)
```

List of 2

```
$ pid: int 787800
$ fd : int [1:2] 6 9
- attr(*, "class")= chr [1:3] "parallelJob" "childProcess" "process"
```

```
1 str(n)
```

List of 2

```
$ pid: int 787801
$ fd : int [1:2] 7 11
- attr(*, "class")= chr [1:3] "parallelJob" "childProcess" "process"
```

mccollect

Checks `mcpallel` objects for completion

```
1 str(mccollect(list(m,n,o)))
```

List of 3

```
$ 787800: num [1:1000000] 1.113 -1.375 0.254 -1.055 0.641 ...  
$ 787801: num [1:1000000] 0.5526 0.0744 0.9768 0.9385 0.1238 ...  
$ 787802: num [1:1000000] 2.00498 4.60403 0.00115 0.58452 0.48502 ...
```

mccollect - waiting

```
1 p = mcpipeline(mean(rnorm(1e5)))
```

```
1 mccollect(p, wait = FALSE, 10)
```

```
$`787803`  
[1] -0.0008566918
```

```
1 mccollect(p, wait = FALSE)
```

```
Warning in selectChildren(jobs, timeout): cannot wait for child  
787803 as it does not exist
```

```
NULL
```

```
1 mccollect(p, wait = FALSE)
```

```
Warning in selectChildren(jobs, timeout): cannot wait for child  
787803 as it does not exist
```

```
NULL
```

doMC & foreach

doMC & foreach

Packages by Revolution Analytics that provides the `foreach` function which is a parallelizable `for` loop (and then some).

- Core functions:
 - `registerDoMC`
 - `foreach, %dopar%, %do%`

registerDoMC

Primarily used to set the number of cores used by `foreach`, by default uses `options("cores")` or half the number of cores found by `detectCores` from the `parallel` package.

```
1 options("cores")
```

```
$cores
```

```
NULL
```

```
1 detectCores()
```

```
[1] 32
```

```
1 getDoParWorkers()
```

```
[1] 1
```

```
1 registerDoMC(4)
2 getDoParWorkers()
```

```
[1] 4
```

foreach

A slightly more powerful version of base `for` loops (think `for` with an `lapply` flavor). Combined with `%do%` or `%dopar%` for single or multicore execution.

```
1 for(i in 1:10) {  
2   sqrt(i)  
3 }
```

```
1 foreach(i = 1:5) %do% {  
2   sqrt(i)  
3 }
```

```
[[1]]
```

```
[1] 1
```

```
[[2]]
```

```
[1] 1.414214
```

```
[[3]]
```

```
[1] 1.732051
```

```
[[4]]
```

```
[1] 2
```

```
[[5]]
```

```
[1] 2.236068
```

foreach - iterators

`foreach` can iterate across more than one value, but it doesn't do length coercion

```
1 foreach(i = 1:5, j = 1:5) %do% {  
2   sqrt(i^2+j^2)  
3 }
```

```
[[1]]  
[1] 1.414214
```

```
[[2]]  
[1] 2.828427
```

```
[[3]]  
[1] 4.242641
```

```
[[4]]  
[1] 5.656854
```

```
[[5]]  
[1] 7.071068
```

```
1 foreach(i = 1:5, j = 1:2) %do% {  
2   sqrt(i^2+j^2)  
3 }
```

```
[[1]]  
[1] 1.414214
```

```
[[2]]  
[1] 2.828427
```

foreach - combining results

```
1 foreach(i = 1:5, .combine='c') %do% {  
2   sqrt(i)  
3 }
```

```
[1] 1.000000 1.414214 1.732051 2.000000 2.236068
```

```
1 foreach(i = 1:5, .combine='cbind') %do% {  
2   sqrt(i)  
3 }
```

```
      result.1 result.2 result.3 result.4 result.5  
[1,]         1 1.414214 1.732051         2 2.236068
```

```
1 foreach(i = 1:5, .combine='+') %do% {  
2   sqrt(i)  
3 }
```

```
[1] 8.382332
```

foreach - parallelization

Swapping out `%do%` for `%dopar%` will use the parallel backend.

```
1 registerDoMC(4)
2 system.time(foreach(i = 1:10) %dopar% mean(rnorm(1e6))))
```

```
user  system elapsed
0.164  0.051   0.090
```

```
1 registerDoMC(8)
2 system.time(foreach(i = 1:10) %dopar% mean(rnorm(1e6))))
```

```
user  system elapsed
0.176  0.098   0.063
```

```
1 registerDoMC(10)
2 system.time(foreach(i = 1:10) %dopar% mean(rnorm(1e6))))
```

```
user  system elapsed
0.202  0.139   0.070
```



furrr / future

```
1 system.time( purrr::map(c(1,1,1), Sys.sleep) )
```

```
   user  system elapsed  
0.000   0.000   3.003
```

```
1 system.time( furrr::future_map(c(1,1,1), Sys.sleep) )
```

```
   user  system elapsed  
0.040   0.003   3.064
```

```
1 future::plan(future::multisession) # See also future::multicore  
2 system.time( furrr::future_map(c(1,1,1), Sys.sleep) )
```

```
   user  system elapsed  
0.392   0.009   1.700
```

Example - Bootstrapping

Bootstrapping is a resampling scheme where the original data is repeatedly reconstructed by taking a samples of size n (with replacement) from the original data, and using that to repeat an analysis procedure of interest. Below is an example of fitting a local regression (`loess`) to some synthetic data, we will construct a bootstrap prediction interval for this model.


```
1 set.seed(3212016)
2 d = data.frame(x = 1:120) |>
3   mutate(y = sin(2*pi*x/120) + runif(length(x),-1,1))
4
5 l = loess(y ~ x, data=d)
6 p = predict(l, se=TRUE)
7
8 d = d |> mutate(
9   pred_y = p$fit,
10  pred_y_se = p$se.fit
11 )
```



```

1  ggplot(d, aes(x,y)) +
2    geom_point(color="gray50") +
3    geom_ribbon(
4      aes(ymin = pred_y - 1.96 * pred_y_se,
5          ymax = pred_y + 1.96 * pred_y_se),
6      fill="red", alpha=0.25
7    ) +
8    geom_line(aes(y=pred_y)) +
9    theme_bw()

```



Bootstrapping Demo

What to use when?

Optimal use of parallelization / multiple cores is hard, there isn't one best solution

- Don't underestimate the overhead cost
- Experimentation is key
- Measure it or it didn't happen
- Be aware of the trade off between developer time and run time

BLAS and LAPACK

Statistics and Linear Algebra

An awful lot of statistics is at its core linear algebra.

For example:

- Linear regression models, find

$$\hat{\beta} = (X^T X)^{-1} X^T y$$

- Principle component analysis
 - Find $T = XW$ where W is a matrix whose columns are the eigenvectors of $X^T X$.
 - Often solved via SVD - Let $X = U\Sigma W^T$ then $T = U\Sigma$.

Numerical Linear Algebra

Not unique to Statistics, these are the type of problems that come up across all areas of numerical computing.

- Numerical linear algebra \neq mathematical linear algebra
- Efficiency and stability of numerical algorithms matter
 - Designing and implementing these algorithms is hard
- Don't reinvent the wheel - common core linear algebra tools (well defined API)

BLAS and LAPACK

Low level algorithms for common linear algebra operations

BLAS

- **B**asic **L**inear **A**lgebra **S**ubprograms
- Copying, scaling, multiplying vectors and matrices
- Origins go back to 1979, written in Fortran

LAPACK

- **L**inear **A**lgebra **P**ackage
- Higher level functionality building on BLAS.
- Linear solvers, eigenvalues, and matrix decompositions
- Origins go back to 1992, mostly Fortran (expanded on LINPACK, EISPACK)

Modern variants?

Most default BLAS and LAPACK implementations (like R's defaults) are somewhat dated

- Written in Fortran and designed for a single cpu core
- Certain (potentially non-optimal) hard coded defaults (e.g. block size).

Multithreaded alternatives:


- ATLAS - Automatically Tuned Linear Algebra Software
- OpenBLAS - fork of GotoBLAS from TACC at UTexas
- Intel MKL - Math Kernel Library, part of Intel's commercial compiler tools
- cuBLAS / Magma - GPU libraries from Nvidia and UTK respectively
- Accelerate / vecLib - Apple's framework for GPU and multicore computing

OpenBLAS Matrix Multiply Performance

```
1 x=matrix(runif(5000^2),ncol=5000)
2
3 sizes = c(100,500,1000,2000,3000,4000,5000)
4 cores = c(1,2,4,8,16)
5
6 sapply(
7   cores,
8   function(n_cores) {
9     flexiblas::flexiblas_set_num_threads(n_cores)
10    sapply(
11      sizes,
12      function(s) {
13        y = x[1:s,1:s]
14        system.time(y %*% y)[3]
15      }
16    )
17  }
18 )
```

n	1 core	2 cores	4 cores	8 cores	16 cores
100	0.000	0.000	0.000	0.000	0.000
500	0.004	0.003	0.002	0.002	0.004
1000	0.028	0.016	0.010	0.007	0.009
2000	0.207	0.110	0.058	0.035	0.039
3000	0.679	0.352	0.183	0.103	0.081
4000	1.587	0.816	0.418	0.227	0.145
5000	3.104	1.583	0.807	0.453	0.266

Matrix Multiply of (n x n) matrices



Matrix Multiply of (n x n) matrices

